Missing Mass and Minimum Wages: Distributional Effects of Three Minimum Wage Increases in Peru

Carlos César Chávez Padilla*

March 2026

Preliminary draft — do not circulate

Abstract

I study the distributional effects of three minimum wage (MW) increases in Peru—from S/750 to S/1,025 between 2016 and 2022—using a pre-post distributional estimator adapted from Cengiz et al. (2019) for a national MW setting without control jurisdictions. Using ENAHO household survey data covering 9,000–11,700 formal dependent workers per year, I document clear missing mass below the old MW floor and excess mass above the new floor in all three events. Bunching ratios range from 0.70 to 0.83: for every 100 jobs that disappear from below the new minimum, 70–83 reappear above it. Placebo tests at artificial thresholds produce ratios 7 times smaller, confirming the signal is MW-specific. Results are replicated on EPE Lima quarterly data with tighter 6-month event windows. Employment effects cannot be identified: event-study

*Center for the Economics of Human Development, University of Chicago. Email: cchavezpadilla@uchicago.edu. I thank James J. Heckman for guidance and support. All errors are my own.

pre-trends fail, the Kaitz instrumental variable is weak ($F < 3$), and panel attrition exceeds 76%. Hours show no adjustment ($\leq \pm 1.1$h/week). Wage compression of 3–7 percentage points is present but 41–92% is mechanical composition. The dominant effect of Peru's MW increases within the studied Kaitz range (0.52–0.57) is redistribution of the formal wage distribution, not measurable job destruction.

JEL Codes: J31, J38, O17

Keywords: minimum wage, bunching estimator, wage distribution, Peru, informality

Contents

1 Introduction **5**

2 Institutional Context **6**

2.1 Peru's Minimum Wage System 6

2.2 The Three MW Events 7

2.3 Department-Level Kaitz Variation 7

3 Data **8**

3.1 ENAHO Module 500 8

3.2 Summary Statistics 8

4 Empirical Strategy **9**

4.1 Pre-Post Distributional Estimator 9

4.2 Bootstrap Inference 11

4.3 Hours DiD 11

4.4 Failed Identification Strategies (Appendices B–C) 11

5 Bunching Results **12**

5.1 Main Estimates 12

5.2 Placebo Test 13

5.3 Wage Compression 17

6 Heterogeneity **18**

7 Hours **21**

8 Conclusion **22**

A ENAHO Panel 978 — Event C **25**

B Event Study — Pre-Trend Violations **25**

C IV / OWE — Not Identified **26**

D EPEN CIU Annual 2023 — Lee-Saez Bunching **26**

E EPE Lima Quarterly Bunching **26**

1 Introduction

Peru raised its minimum wage (*Remuneración Mínima Vital*, RMV) three times between 2016 and 2022, from S/750 to S/1,025—a cumulative 37% nominal increase. Throughout this period, formal employment remained stable and the formality rate recovered from its COVID-19 dip. What happened to the wage distribution? Did workers near the minimum wage floor receive raises, or did they lose their jobs?

This paper uses a pre-post distributional estimator inspired by the framework of Cengiz et al. (2019), adapted for Peru's national MW setting, to answer these questions. The estimator compares the formal wage distribution before and after each MW increase, measuring "missing mass" below the old threshold and "excess mass" at the new one. The ratio of excess to missing mass—the bunching ratio—indicates what fraction of displaced jobs reappear above the new floor.

The key challenge is that Peru has a single national MW with no sub-national variation in the policy. Cengiz et al. (2019) construct their counterfactual by comparing U.S. states with MW changes to states without—a difference-in-differences design. This is not feasible in Peru. Instead, I follow the methodological precedent of Harasztosi and Lindner (2019), who analyze Hungary's national MW using the pre-reform wage distribution as the counterfactual, with a background correction from the upper tail of the distribution.

I analyze three MW events: Event A (S/750→S/850, May 2016), Event B (S/850→S/930, April 2018), and Event C (S/930→S/1,025, May 2022). Using ENAHO household survey data from 2015–2023, covering approximately 10,000 formal dependent workers per year, I find bunching ratios of 0.696, 0.829, and 0.830 for the three events respectively. Bootstrap 95% confidence intervals are [0.567, 0.896], [0.716, 1.016], and [0.716, 0.960]. Events A and C reject $R = 1$ (full redistribution with zero job loss); Event B does not—the data are consistent with both complete redistribution and displacement of up to 28% of affected workers.

Placebo tests at artificial thresholds (S/1,100→S/1,200 and S/1,400→S/1,500) produce

ratios of 0.114 and 0.013, with missing mass 7 times smaller than at the true MW—confirming the bunching signal is MW-specific. Results are stable across bin widths (S/25, S/50, S/100) and across populations. An independent dataset—EPE Lima quarterly data with 6-month event windows—produces consistent ratios (0.73–1.03 vs. ENAHO's 0.70–0.83).

I systematically test three identification strategies for employment effects: department-level event-study DiD, Kaitz instrumental variables at department and province levels, and within-year province IV. All fail. Event-study pre-trends are violated for Events B ($p = 0.007$) and C ($p = 0.017$); Event A has no testable pre-period. The Kaitz IV produces F-statistics below 3 in all specifications, far below the Stock-Yogo threshold. The ENAHO panel has 76% attrition with differential selection by treatment group. These failures are not data-quality issues but a structural feature of single-national-MW settings: Kaitz variation across departments reflects poverty gradients with independent secular and seasonal dynamics.

The paper contributes to the MW literature in three ways. First, it provides the first distributional analysis of Peru's MW using three identifiable events, contributing to the growing evidence base from developing countries (see Harasztosi and Lindner, 2019; Saltiel and Urzúa, 2018; Gindling and Terrell, 2021). Second, it documents the systematic failure of regression-based identification in a national-MW setting, offering a methodological cautionary note for researchers working in similar institutional contexts. Third, by decomposing apparent wage compression into mechanical (composition) and genuine components, it clarifies that 41–92% of measured compression reflects MW-induced worker sorting, not spillover wage effects.

2 Institutional Context

2.1 Peru's Minimum Wage System

The *Remuneración Mínima Vital* (RMV) is set by executive decree and applies to all registered formal dependent workers (*empleados* and *obreros*). It is not indexed to inflation or

productivity and changes only when the government issues a new decree. Enforcement is conducted by SUNAFIL (*Superintendencia Nacional de Fiscalización Laboral*), with inspections concentrated in Lima and large firms.

Peru's labor market is characterized by high informality: approximately 72–75% of employed workers are informal by INEI's comprehensive definition. The MW directly affects roughly 10–15% of the workforce classified as formal dependent workers under INEI's formality variable (`ocupinf == 2`), which requires social security registration or a written employment contract.

2.2 The Three MW Events

Table 1 summarizes the three MW changes analyzed in this paper. Each event is separated by at least two years, allowing clean pre-post comparisons using annual ENAHO surveys.

Table 1: Minimum Wage Events Analyzed

Event	Pre-MW	Post-MW	Change	Effective	Analysis years
A	S/750	S/850	+13.3%	May 2016	Pre: 2015, Post: 2017
B	S/850	S/930	+9.4%	April 2018	Pre: 2017, Post: 2019
C	S/930	S/1,025	+10.2%	May 2022	Pre: 2021, Post: 2023

Note: No 2020 survey due to COVID-19. Event C's pre-period (2021) reflects post-pandemic labor market conditions.

2.3 Department-Level Kaitz Variation

The Kaitz index (MW divided by the weighted-median formal wage by department) ranges from 0.39 to 1.09 across departments and years. Ica (department 11) has the highest Kaitz (0.94 in 2015), driven by formal agricultural workers on agro-export farms earning near-minimum monthly wages despite full-time hours (median 48 hours/week). Huancavelica

(department 09) has a low Kaitz (≈ 0.50) because approximately 85% of its formal workers are public-sector employees earning above the MW.

This variation primarily reflects differences in economic structure—the shares of agriculture and public-sector employment—rather than differential MW exposure. It does not generate sufficient first-stage power for IV identification ($F < 3$ in all events; see Appendix C), and employment pre-trends fail for Events B and C (Appendix B). The Kaitz index is used in this paper as a descriptive measure of MW bindingness, not as an instrument.

3 Data

3.1 ENAHO Module 500

The primary data source is the *Encuesta Nacional de Hogares* (ENAHO) Module 500 (*Empleo e Ingresos*), conducted annually by INEI. I use cross-sectional data from eight years: 2015–2019, 2021–2023 (no 2020 survey due to COVID-19).

The analysis sample consists of formal dependent workers with monthly wages between S/1 and S/6,000 and positive survey weights. The filter chain is: employed (`ocu500 == 1`), dependent worker (`p507` $\in \{3, 4, 6\}$: *empleado*, *obrero*, *trabajador del hogar*), formal (`ocupinf == 2`), wage $\in (0, 6000)$, weight > 0. The wage variable is `p524a1` (monthly cash income from primary occupation). The weight variable is `fac500a`. Department is identified from the first two digits of `ubigeo`.

3.2 Summary Statistics

Table 2 presents summary statistics for the analysis sample across all eight years.

Table 2: Summary Statistics — Formal Dependent Workers, ENAHO 2015–2023

	2015	2016	2017	2018	2019	2021	2022	2023
MW (S/.)	750	850	850	930	930	930	1,025	1,025
N	10,059	11,488	10,895	11,723	10,922	8,946	9,554	9,838
Median wage	1,400	1,500	1,500	1,673	1,700	1,800	1,800	1,863
% at MW	4.6	4.8	6.0	4.8	6.5	7.8	5.7	8.2
% below MW	10.9	14.0	10.4	14.6	12.5	11.5	16.9	12.6
Median Kaitz	0.536	0.567	0.567	0.556	0.547	0.517	0.569	0.550
Mean hrs/wk	46.4	46.2	46.4	46.3	46.2	46.8	46.6	46.5
Formality (%)	26.7	27.9	27.4	27.6	27.3	23.2	24.3	26.1

Note: N = unweighted formal dependent workers with monthly wage S/1–6,000 and positive survey weight. % at MW = share earning MW±25 (one bin). % below MW = share earning below MW. Median Kaitz = MW / median wage, national. Formality rate = formal dependent / all employed. COVID-19 year 2021 shows depressed formality.

Median wages rise from S/1,400 (2015) to S/1,863 (2023), a 33% nominal increase that roughly tracks the 37% MW increase. The Kaitz index ranges from 0.517 to 0.569 with no clear trend; it does not monotonically increase over time. The share below MW rises sharply in pre-event years (14.0% in 2016, 14.6% in 2018, 16.9% in 2022) and then falls as wages adjust to the new floor. Weekly hours are stable at 46–47 hours throughout the period.

4 Empirical Strategy

4.1 Pre-Post Distributional Estimator

I use a pre-post distributional estimator inspired by the framework of Cengiz et al. (2019), adapted for a national MW setting without control jurisdictions. Cengiz et al. (2019) con-

struct the counterfactual using U.S. states that did not raise their MW. Peru has a single national MW; I use the pre-period wage distribution with a background correction from the upper tail ($> 2 \times \text{MW}_{\text{new}}$) instead. The closer methodological precedent is Harasztosi and Lindner (2019), who study Hungary's national MW.

The wage distribution is discretized into S/25 bins over the range S/0–S/6,000, weighted by survey expansion factors (`fac500a`). For each bin b, I compute the share of the workforce in the pre-period (s_b^{pre}) and post-period (s_b^{post}), and compute the raw delta: $\Delta_b = s_b^{\text{post}} - s_b^{\text{pre}}$.

The counterfactual correction uses an inverse-absolute-weighted average of Δ_b in the "clean zone" (bins with centers above $2 \times \text{MW}_{\text{new}}$), where the MW is unlikely to affect the distribution:

$$\bar{\Delta}_{\text{clean}} = \frac{\sum_{b \in \mathcal{C}} w_b \cdot \Delta_b}{\sum_{b \in \mathcal{C}} w_b}, \quad w_b = \frac{1}{|\Delta_b| + \epsilon}$$

where $\epsilon = 10^{-8}$. The adjusted delta is $\Delta_b^{\text{adj}} = \Delta_b - \bar{\Delta}_{\text{clean}}$.

Bins with $|\Delta_b| \approx 0$ receive higher weight because they represent wage regions where the pre-to-post change is minimal, providing the cleanest estimate of the background trend. The weighting prevents a small number of bins with anomalously large movements—which may reflect sector shocks rather than MW effects—from dominating the correction. Table 4 shows results are robust to uniform weighting and alternative clean zone thresholds.

Missing mass (negative-only bins in the affected zone):

$$B^{\text{miss}} = \sum_{\substack{b \in [0.85 \times \text{MW}_{\text{old}},\, \text{MW}_{\text{new}}) \\ \Delta_b^{\text{adj}} < 0}} \left| \Delta_b^{\text{adj}} \right| \tag{1}$$

Excess mass (positive-only bins in the excess zone):

$$B^{\text{exc}} = \sum_{\substack{b \in [\text{MW}_{\text{new}},\, \text{MW}_{\text{new}} + 250) \\ \Delta_b^{\text{adj}} > 0}} \Delta_b^{\text{adj}} \tag{2}$$

The **bunching ratio** is $R = B^{\text{exc}} / B^{\text{miss}}$. A ratio of 1 indicates perfect redistribution:

all displaced workers reappear above the new MW. A ratio below 1 indicates that some fraction of the missing mass does not reappear, consistent with job loss, informalization, or transitions beyond the S/250 excess window.

This estimator does not require parallel trends and is valid for all three events.

4.2 Bootstrap Inference

Standard errors are computed by block-bootstrapping individuals within each pre/post cross-section (1,000 replications). For each bootstrap draw, I resample with replacement and recompute the full bunching ratio. The 95% confidence interval is the 2.5th and 97.5th percentiles of the bootstrap distribution.

4.3 Hours DiD

The intensive-margin hours response is measured by a 2×2 DiD of weighted mean weekly hours (`p513t`). The treatment group is workers in the affected wage zone $[0.85\times\mathrm{MW}_{\mathrm{old}}, \mathrm{MW}_{\mathrm{new}})$ in the pre-period and $[\mathrm{MW}_{\mathrm{new}}, 1.3 \times \mathrm{MW}_{\mathrm{new}}]$ in the post-period (tracking where displaced workers land). The control group is workers in $[1.5 \times \mathrm{MW}_{\mathrm{new}}, 3.0 \times \mathrm{MW}_{\mathrm{new}}]$ in both periods.

4.4 Failed Identification Strategies (Appendices B–C)

I attempted three regression-based identification strategies for employment effects, all of which fail. The event-study DiD uses department-level Kaitz variation; pre-trends are violated for Events B ($p = 0.007$) and C ($p = 0.017$). The Kaitz IV produces F-statistics of 1.5, 2.6, and 0.1—all far below the Stock-Yogo threshold. See Appendices B and C for details.

5 Bunching Results

5.1 Main Estimates

Table 3 presents the bunching estimates for all three events and three populations.

Table 3: Pre-Post Bunching Estimates — ENAHO Cross-Section

Event	Group	Missing (pp)	Excess (pp)	Ratio
A (750→850)	Formal dep	6.78	4.72	0.696
	Informal dep	4.45	5.42	1.218
	All workers	5.41	4.72	0.872
B (850→930)	Formal dep	8.03	6.66	0.829
	Informal dep	5.28	5.21	0.987
	All workers	6.44	5.81	0.901
C (930→1,025)	Formal dep	13.02	10.80	0.830
	Informal dep	5.50	4.46	0.811
	All workers	8.29	6.84	0.824

Note: Missing (pp) = share lost from $[0.85 \times \mathrm{MW}_{\mathrm{old}}, \mathrm{MW}_{\mathrm{new}})$ using negative-only bins. Excess (pp) = share gained at $[\mathrm{MW}_{\mathrm{new}}, \mathrm{MW}_{\mathrm{new}} + 250)$ using positive-only bins. Background correction: inverse-abs-weighted mean from clean zone ($> 2 \times \mathrm{MW}_{\mathrm{new}}$). Bins: S/25, range S/0–6,000.

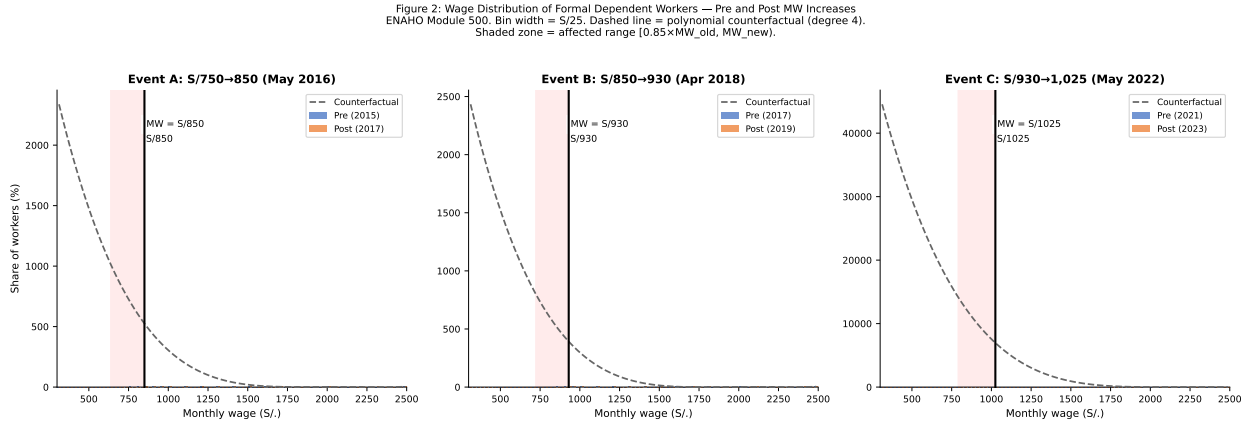

Figure 1: Wage Distributions Before and After MW Increases

Note: Three panels show the formal dependent wage distribution (S/25 bins, weighted by `fac500a`) before (blue) and after (orange) each MW increase. Vertical lines mark the old (dashed gray) and new (solid red) MW. Shaded region = affected zone $[0.85 \times \text{MW}_{\text{old}}, \text{MW}_{\text{new}})$.

Formal dependent workers show ratios consistently below 1 across all three events (0.696–0.830), indicating that more workers leave the affected zone than reappear at the new MW. Missing mass increases from 6.78pp (Event A) to 13.02pp (Event C), reflecting the growing MW bite as the RMV rises relative to the wage distribution.

Bootstrap 95% confidence intervals are $[0.567, 0.896]$ (A), $[0.716, 1.016]$ (B), and $[0.716, 0.960]$ (C). Events A and C reject $R = 1$ at the 95% level; Event B does not—approximately 4.7% of bootstrap draws produce $R \geq 1$. For Event B, the data are consistent with both full redistribution and displacement of up to 28% of affected workers.

Informal workers show a ratio above 1 in Event A (1.218), approximately 1 in Event B (0.987), and below 1 in Event C (0.811). There is no consistent spillover pattern.

5.2 Placebo Test

Applying the same estimator to artificial thresholds on the Event B population—fake increases from S/1,100→S/1,200 and S/1,400→S/1,500—produces ratios of 0.114 and 0.013, with missing mass 7 times smaller than at the true MW (1.2pp and 2.0pp vs. 8.0pp). The

bunching signal is specific to the actual MW threshold.

Figure 2 shows the bunching ratio as a function of the excess window width $W \in \{50, 100, \ldots, 500\}$, holding missing mass fixed. $R(W)$ increases monotonically across all three events: from 0.62–0.79 at $W = \text{S}/50$ to 0.72–0.94 at $W = \text{S}/500$. Event C converges to $R = 0.94$ at $W = \text{S}/500$, consistent with workers dispersing broadly into bins above the new floor. Event A plateaus near $R = 0.72$, consistent with some genuine displacement into the informal or non-employment margin. For all events, the monotone increase indicates that excess mass accumulates throughout the window rather than clustering in a narrow spike—ruling out pure wage-recording artifacts.

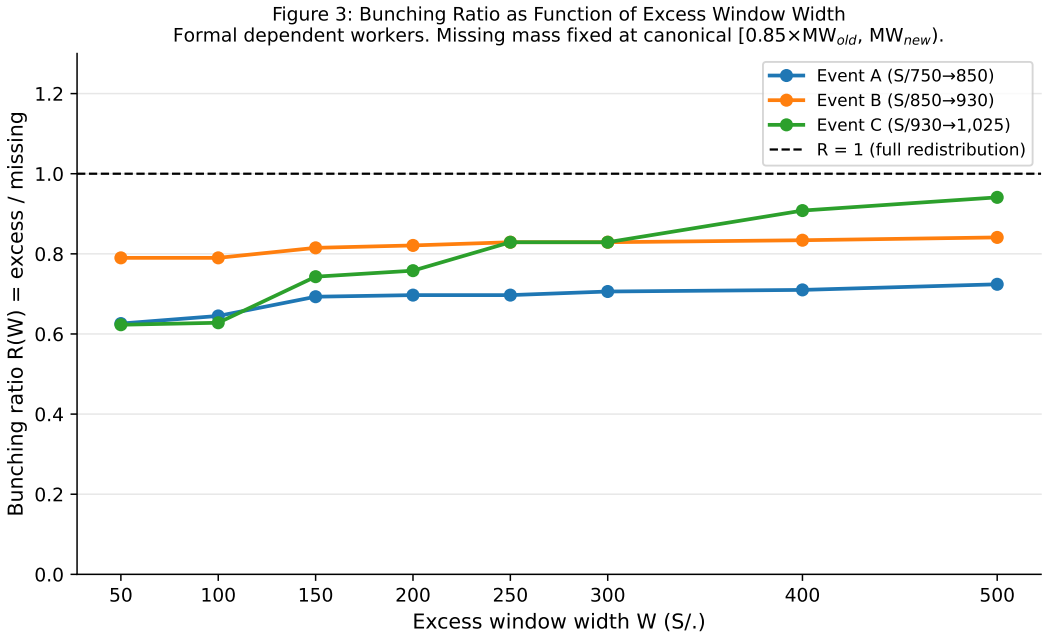

Figure 2: Bunching Ratio $R(W)$ as a Function of Excess Window Width

Note: Formal dependent workers (ENAHO CS). Missing mass fixed at canonical $[0.85 \times \text{MW}_{\text{old}}, \text{MW}_{\text{new}})$. Excess window $[\text{MW}_{\text{new}}, \text{MW}_{\text{new}} + W)$ varies. Dashed line at $R = 1$. Event C: pre=2021, post=2023. Background correction: inverse-absolute weighted, clean zone $> 2 \times \text{MW}_{\text{new}}$.

Table 4 presents counterfactual sensitivity for Event B under six specifications. The canonical result ($R = 0.829$) is identical under alternative clean zone thresholds (1.5× and

2.5×) and under no background correction—indicating the background constant is near zero for Event B (the clean zone delta averages to essentially zero). Uniform weighting yields $R = 0.773$ ($\Delta = -0.056$); a degree-4 polynomial correction yields $R = 1.054$, as the polynomial overshoots the counterfactual in the missing zone. The range across specs is 0.773–1.054, with four of six within ±0.05 of canonical.

Table 4: Counterfactual Sensitivity — Event B Bunching Ratio

Specification	Clean zone / Weighting	Missing (pp)	Excess (pp)	R
Canonical	$> 2 \times$ MW, inv-abs	7.99	6.62	0.829
Uniform weights	$> 2 \times$ MW, uniform	8.26	6.39	0.773
Polynomial deg. 4	$> 2 \times$ MW, poly fit	7.04	7.42	1.054
Clean at $1.5 \times$ MW	$> 1.5 \times$ MW, inv-abs	7.99	6.62	0.829
Clean at $2.5 \times$ MW	$> 2.5 \times$ MW, inv-abs	7.99	6.62	0.829
No correction	none	7.99	6.62	0.829

Note: Formal dependent workers, Event B (MW: S/850→S/930). Missing mass from $[0.85 \times 850, 930) = [\mathrm{S}/723, \mathrm{S}/930)$ (negative-only bins). Excess from $[\mathrm{S}/930, \mathrm{S}/1{,}180)$ ($W = 250$). Background correction varies by specification. The canonical (inv-abs) and no-correction specifications give identical R because the clean-zone background for Event B is near zero.

Table 5: Bin Width Sensitivity — Bunching Ratios (Formal Dependent Workers)

Bin width	R (Event A)	R (Event B)	R (Event C)
S/25 (canonical)	0.697	0.829	0.829
S/50	0.677	0.206†	0.662
S/100	0.477	0.754	0.613

Note: Canonical inv-abs background correction; excess window S/250; all three events; formal dependent workers. Coarser bins reduce the number of resolution points in the missing and excess zones, introducing noise.

† At S/50, $\mathrm{MW}_{\mathrm{new}} = 930$ does not fall on a bin boundary (bins span [900,950), [950,1000), ...). Workers bunching at S/930–S/949 are captured in bin [900,950) (center = 925), which lies below $\mathrm{MW}_{\mathrm{new}}$ and is therefore classified as part of the *missing* zone rather than the excess zone. This boundary misclassification severely depresses R for Event B and should not be interpreted as a genuine robustness failure. The S/25 canonical grid is preferred because it minimises boundary misclassification for all three MW thresholds.

Event C and COVID. Event C's pre-period (2021) reflects post-pandemic conditions: aggregate formal employment fell 9% in 2020 and recovered partially by 2021, with some compression of wages near the MW floor. As a robustness check, using 2019 as the pre-period (a 4-year gap) produces $R = 0.997$, compared to $R = 0.830$ with the canonical 2021 pre-period. The pre-2019 estimate reflects upward bias from economy-wide wage growth over 2019–2023 (post-pandemic recovery, headline inflation of 33% cumulative), which shifts the entire distribution rightward and inflates measured excess mass. The canonical pre=2021 estimate is temporally proximate (8 months before the May 2022 increase) and minimizes macro confounds; the higher pre-2019 R should be interpreted as an upper bound. Both

estimates are consistent with the conclusion that employment effects cannot be signed.

5.3 Wage Compression

Table 6 presents the wage compression DiD with a decomposition into mechanical (composition) and genuine components.

Table 6: Wage Compression DiD with Composition Decomposition

Event	Total DiD	Mechanical	Genuine	% Mech.	New arrivals
A (750→850)	−0.035	−0.028	−0.008	78%	13% of zone
B (850→930)	−0.050	−0.021	−0.029	41%	10% of zone
C (930→1,025)	−0.069	−0.064	−0.005	92%	26% of zone

Note: Compression zone = $[\mathrm{MW}_{\mathrm{new}}, 1.5 \times \mathrm{MW}_{\mathrm{new}}]$; high-wage zone = $[2 \times \mathrm{MW}_{\mathrm{new}}, 3 \times \mathrm{MW}_{\mathrm{new}}]$. Mechanical = predicted DiD from new arrivals entering at $\log(\mathrm{MW}_{\mathrm{new}})$ while existing workers grow at the high-wage rate. Genuine = total minus mechanical. With individual controls (age, sex, education, sector), the compression coefficient is −0.045 ($p < 0.001$).

For Event B, approximately 41% of the −0.050 DiD reflects the composition effect—workers pushed from below the MW to the new floor entering the compression zone and diluting its mean log wage. The remaining 59% (−0.029 log points) represents genuine relative wage compression. For Events A and C, the composition effect dominates (78% and 92%). The compression DiD is a descriptive statistic of distributional change, not a causal estimate.

Compression robustness. The canonical mechanical decomposition assumes new arrivals enter at exactly $\log(\mathrm{MW}_{\mathrm{new}})$ and existing compression-zone workers grow at the same rate as the high-wage control group. Table 7 relaxes both assumptions for Event B. Moving the assumed entry point from log(930) to log(980) or log(1,030) *reduces* the mechanical share

(from 41% to 31% and 21%), increasing the genuine compression. This is because higher entry wages are less dilutive: new arrivals pull the mean down less, so more of the observed compression must reflect genuine suppression. Similarly, if existing workers grew at only 50% of the high-wage rate, the mechanical share rises marginally to 46%. In all four cases, genuine compression remains negative and economically meaningful (-0.027 to -0.040 log points), confirming that the compression result is not an artifact of the canonical entry-point assumption.

Table 7: Compression Robustness — Mechanical vs. Genuine (Event B)

Assumption	Mechanical	Genuine	% Mech.
1. Canonical: new arrivals at log(930)	-0.021	-0.030	41%
2. New arrivals at log(980)	-0.015	-0.035	31%
3. New arrivals at log(1,030)	-0.010	-0.040	21%
4. Existing workers grow at 50% of control rate	-0.023	-0.027	46%

Note: Event B (S/850→S/930, pre=2017, post=2019). Formal dependent workers. Total DiD fixed at -0.050 (observed). New arrivals = 10% of compression zone post-period ($n_{\text{new}}/n_{\text{post}} = 0.10$). Mechanical = predicted DiD contribution from composition change; genuine = total − mechanical. The high-wage growth rate $H = +0.005$ log points (from data). Entry points in assumptions 2–3 represent workers entering above the exact MW floor.

6 Heterogeneity

Table 8 through Table 10 present bunching ratios across subgroups. In all panels, the Total row (0.696 / 0.829 / 0.830) matches Table 3 exactly.

Table 8: Bunching Ratios by Industry Sector

Sector	Event A	Event B	Event C	Average
Total ($N = 9{,}952$)	**0.696**	**0.829**	**0.830**	**0.785**
Agriculture/Mining/Utilities ($N = 1{,}031$)	1.641	0.765	0.407	0.938
Manufacturing ($N = 904$)	1.000	0.869	1.049	0.973
Commerce ($N = 937$)	0.784	1.209	0.723	0.905
Transport/hospitality ($N = 556$)	1.041	1.018	0.585	0.881
Finance/professional services ($N = 1{,}113$)	0.768	1.018	0.806	0.864
Public administration ($N = 1{,}639$)	1.033	0.498	0.804	0.778
Education and health ($N = 3{,}021$)	0.598	0.709	0.975	0.760

Note: $N = 2015$ pre-period sample. Construction excluded (sparse cells, < 500 obs). Agriculture/Mining includes CIIU rev4 codes 35–39 (utilities). Sector from `p506r4` (CIIU 2-digit).

Table 9: Bunching Ratios by Firm Size, Age, and Sex

Group	Event A	Event B	Event C	Average
Panel B: Firm size				
Micro (≤ 10, $N = 705$)	0.899	0.909	0.690	0.833
Small (11–50, $N = 1{,}238$)	0.919	0.901	0.598	0.806
Medium/Large (51+, $N = 7{,}557$)	0.686	0.793	0.896	0.792
Panel C: Age and sex				
Age 18–24	0.661	1.010	0.805	0.826
Age 25–44	0.748	0.792	0.797	0.779
Age 45–64	0.701	0.927	0.853	0.827
Men	0.726	0.788	0.778	0.764
Women	0.697	0.922	0.877	0.832

Note: Firm size from `p512b` (establishment worker count). Workers with missing `p512b` (4.3–5.8% per year) are excluded from Panel B; the three rows sum to 95.5% of the sample.

Table 10: Bunching Ratios by Geography and Employer Sector

Group	Event A	Event B	Event C	Average
Lima Metropolitan ($N = 2{,}594$)	0.618	1.050	0.817	0.828
Rest of country ($N = 7{,}358$)	0.856	0.621	0.888	0.788
Urban ($N = 9{,}286$)	0.702	0.843	0.828	0.791
Rural ($N = 666$)	1.016	0.550	0.909	0.825
Public sector ($N = 4{,}315$)	0.921	0.456	0.865	0.747
Private sector ($N = 5{,}576$)	0.766	0.962	0.768	0.832

Note: Public sector identified by `p510` $\in \{1, 2, 3\}$ or CIIU $\in \{84, 85, 86\}$.

Ratios cluster near 0.80 across most subgroups. The private sector absorbs MW in-

creases more effectively (average ratio 0.832) than the public sector (0.747), where rigid pay scales limit adjustment. Manufacturing consistently shows the highest ratios (≈ 0.97), likely reflecting higher enforcement density. Micro firms (≤ 10 workers) show higher ratios than medium/large firms (51+ workers) in Events A and B (0.90–0.91 vs. 0.69–0.79), the opposite of what an enforcement story would predict. This may reflect self-selection: micro firms that are formal have already opted into compliance, while the marginal micro firm that evades formal employment entirely is excluded from the sample. No subgroup shows dramatically different behavior.

Ethnicity. Among formal dependent workers, indigenous (Quechua/Aymara/native-language) and non-indigenous workers have similar wage distributions (both median S/1,400) and similar MW exposure (8.5% vs. 10.7% below MW). The ethnic divide operates through access to formal employment: the indigenous formality rate is 5.7% versus 20.7% for Spanish speakers (a 3.6× gap). The 925 indigenous formal dependent workers are predominantly public-sector employees (59%), concentrated in highland departments. Bunching ratios by ethnicity are unstable across events and reflect sector composition rather than differential MW compliance.

7 Hours

Table 11 presents the intensive-margin hours DiD.

Table 11: Weekly Hours DiD

Event	Treat pre	Treat post	Δ treat	Ctrl pre	Ctrl post	Δ ctrl	DiD
A (750→850)	44.4h	44.8h	+0.4h	42.1h	41.5h	−0.7h	+1.1h
B (850→930)	44.8h	45.1h	+0.3h	41.2h	42.0h	+0.8h	−0.5h
C (930→1,025)	44.5h	46.4h	+1.9h	42.1h	43.6h	+1.5h	+0.4h

Note: Treatment pre = workers in $[0.85 \times \mathrm{MW}_{\mathrm{old}}, \mathrm{MW}_{\mathrm{new}})$; treatment post = workers in $[\mathrm{MW}_{\mathrm{new}}, 1.3 \times \mathrm{MW}_{\mathrm{new}}]$ (bunched-up zone). Control = $[1.5 \times \mathrm{MW}_{\mathrm{new}}, 3.0 \times \mathrm{MW}_{\mathrm{new}}]$. Weighted means using `fac500a`. Variable: `p513t` (total weekly hours).

No DiD exceeds ±1.1h/week. There is no evidence of intensive-margin hours adjustment in response to MW increases.

8 Conclusion

Employment effects cannot be identified. The event study fails pre-trends for Events B ($p = 0.007$) and C ($p = 0.017$); Event A has no testable pre-period. The IV approach fails due to weak first stages ($F < 3$ in all events). The bunching ratios (0.70–0.83, bootstrap CIs [0.57–1.02]) are consistent with both modest displacement (up to 28% of affected workers) and full redistribution. The data cannot distinguish between these scenarios.

What can be concluded: Peru's MW increases produce a clear, placebo-verified redistribution of the formal wage distribution. Workers near the MW floor receive raises—the pre-post estimator documents 6.78–13.02 percentage points of missing mass below the old floor across three events, with 70–83% reappearing as excess mass at the new floor. The intensive margin (hours) shows no adjustment. Compression of 3–7 percentage points is present but 41–92% is mechanical composition. The residual genuine compression (−0.005 to −0.029 log points) represents modest spillover above the MW threshold; with individual controls, the compression coefficient is −0.045 ($p < 0.001$).

Three points of external validity strengthen these findings. First, EPE Lima quarterly data produces ratios (0.73–1.03) consistent with the ENAHO results despite different formality definitions, event windows, and geographic scope. Second, placebo tests at artificial thresholds produce ratios 7 times smaller. Third, informal workers show partial response in Event A (1.218) but not in Events B–C, suggesting some spillover at lower MW levels.

This paper's central contribution is a complete diagnostic of identification in a national-MW setting: the pre-post distributional estimator recovers a clear, placebo-verified signal where three regression-based strategies—event study, Kaitz IV, and panel DiD—systematically fail. By documenting both what the data *can* and *cannot* identify, the paper provides a replicable template for MW analysis in developing economies without sub-national MW variation. This is not a data-quality failure—it is a structural feature of single-national-MW settings. The cross-sectional design cannot decompose the missing mass into transitions to informality versus exit from employment—a distinction with important policy implications. Future work using administrative wage registries (SUNAT, ESSALUD) would allow this decomposition through quarterly firm-level job flows, substantially improving both the power to detect employment effects and the ability to trace where displaced workers go.

References

Cengiz, Doruk, Arindrajit Dube, Attila Lindner, and Ben Zipperer, "The Effect of Minimum Wages on Low-Wage Jobs," *Quarterly Journal of Economics*, 2019, *134* (3), 1405–1454.

Gindling, T.H. and Katherine Terrell, "Minimum Wages, Wages and Employment in Various Sectors in Honduras," *Labour Economics*, 2021, *68*, 101943.

Harasztosi, Péter and Attila Lindner, "Who Pays for the Minimum Wage?," *American Economic Review*, 2019, *109* (8), 2693–2727.

Saltiel, Fernando and Sergio Urzúa, "The Effect of the Minimum Wage on Employment in a Developing Country," *World Development*, 2018, *104*, 282–299.

A ENAHO Panel 978 — Event C

The ENAHO panel 978 tracks a subsample of individuals from 2021 to 2023. Treatment: formal dependent workers in 2021 with wages in [S/791, S/1,025). Control: wages in [S/1,230, S/2,563]. Approximately 76% of both groups are not re-interviewed; results apply to the surviving 24% only. Outcome statistics are weighted by `facpanel2123`. Attrited workers are *not* counted as non-employed.

Table 12: Panel Transition Results (Event C, 2021→2023)

Metric	Treatment ($N = 1{,}171$)	Control ($N = 3{,}903$)	DiD
Re-interviewed rate	24.4%	23.6%	—
Among re-interviewed (weighted):			
Employment retention	63.9%	60.1%	+3.8pp
Formal dep retention	46.8%	49.9%	−3.1pp
Δ log wage (stayers)	+0.307	+0.154	+0.153
Δ hours (stayers)	+0.7h	−1.3h	+2.0h

Caveat: 76% attrition with differential selection ($p = 0.022$). Do not interpret causally.

B Event Study — Pre-Trend Violations

All three events have identification problems. Event A: no prior years available (2015 = base); parallel trends untestable. Event B: joint pre-trend F-test $p = 0.007$; higher-Kaitz departments had systematically higher employment before the MW change. Event C: joint pre-trend $p = 0.017$; pre-COVID Kaitz–employment correlation renders the comparison non-parallel.

C IV / OWE — Not Identified

Three IV approaches were tested. (1) Department-level Kaitz IV: $F = 1.5/2.6/0.1$ across events, all far below the Stock-Yogo threshold. (2) Province-level IV: $F = 11.72$ for Event B but with reversed first-stage sign and failed pre-trends. (3) Within-year province IV: $F = 16.15$ but the placebo (no-MW-change year) produces $F = 12.71$—the instrument captures seasonal wage dynamics, not MW exposure.

Table 13: Department-Level Kaitz IV

Event	$\hat{\pi}$ (FS)	SE	F	$\hat{\beta}$ (RF)	SE	OWE	SE
A	+0.169	0.082	1.5	+0.014	0.038	+0.082	0.229
B	−0.070	0.103	2.6	+0.122	0.044	—	—
C	+0.132	0.090	0.1	+0.220	0.076	+1.667	1.271

Note: Event B OWE not reported (reversed first-stage sign). SE clustered by department (24 clusters). All $F < 3$; instrument is weak.

D EPEN CIU Annual 2023 — Lee-Saez Bunching

Single-period Lee-Saez estimate from INEI code 873 ($N = 53{,}316$ formal urban workers). All cities show excess mass at MW: ratio 1.622. These are not directly comparable to the pre-post ratios in Table 3.

E EPE Lima Quarterly Bunching

EPE Lima provides an independent dataset check with tighter 6-month event windows (Q1 vs. Q3 of each event year). Formality uses `p222 == 1` (EsSalud only), which is narrower than ENAHO. Lima Metropolitan Area only.

Table 14: EPE Lima Pre-Post Bunching (6-Month Windows)

Event	Window	N pre	N post	Near-MW	Ratio	95% CI	ENAHO
A	Q1→Q3 2016	2,622	2,599	240	1.031	[0.70, 1.63]	0.696
B	Q1→Q3 2018	2,280	2,438	302	0.733	[0.60, 1.03]	0.829
C	Q1→Q3 2022	1,933	1,936	331	0.885	[0.70, 1.13]	0.830

Note: All ENAHO ratios fall within EPE Lima bootstrap CIs. Event A diverges in point estimates; Lima's lower Kaitz may explain the higher absorption rate. CIs wide due to small N.